

Pattern Classification (EET 3035)

Lecture 01

Dr. Kundan Kumar
PhD (IIT Kharagpur)
Associate Professor
Department of ECE

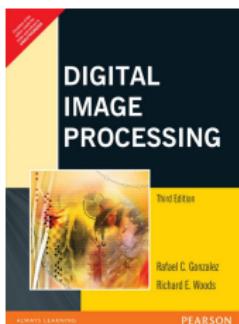
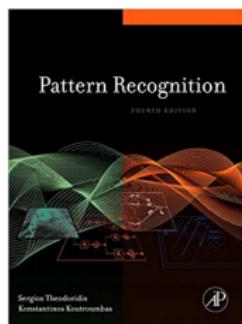
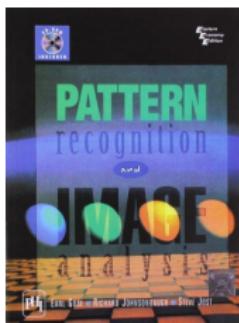
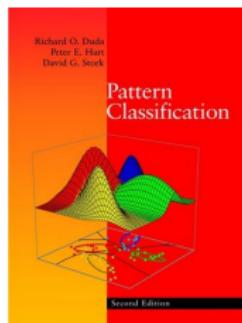


Faculty of Engineering (ITER)
S'O'A Deemed to be University, Bhubaneswar, India-751030
© 2020 Kundan Kumar, All Rights Reserved

Outline

- 1 Text books and Syllabus
- 2 Introduction
- 3 Applications
- 4 An Example
- 5 Pattern Recognition System
- 6 Design Cycle
- 7 References

Text Books and Credits



Text Book:

- Pattern Classification, Duda-Hart, 2nd Edition (**Prescribed**)
 - Pattern Recognition and Image Analysis by Earl Gose
 - Pattern Recognition by Theodoridis, 4th Edition
 - Digital Image Processing by Gonzalez, 3rd Edition

Credits:

- 4 credits course, 4 Classes/week (1hr/Class)
 - Prerequisite: MTH 2002 (Probability and Statistics)

Grading Pattern and Syllabus

■ **Grading pattern: 6**

Attendance	:	5 Marks
2 Quizzes	:	10 Marks
Assignments	:	10 Marks
Mid-term examination	:	15 Marks
Total Internal	:	40 Marks

Theory examination	:	60 Marks
Total External	:	60 Marks

■ **Syllabus:** Introduction, Features, Bayesian Decision Theory, Maximum Likelihood and Bayesian Parameter Estimation, Non-parameteric Techniques, Linear Discriminant Functions, Multilayer Neural Networks, Unsupervised learning and Clustering

Google Classrooms

- All the communication will be through Google Classroom:
 - Course materials
 - Assignments
 - Announcements and Notices
- join the course at <https://classroom.google.com/>
- Class Code for section C:

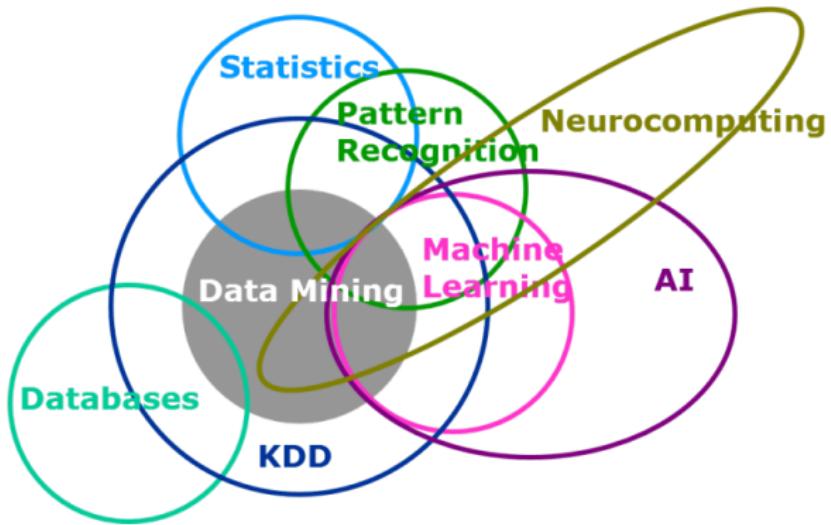
kkne5zl

- Class Code for section D:

neqep3k

- Would you like to code in Python for the topics to be covered in Pattern Classification?

Introduction to Pattern Classification



Relation with AI and ML

■ Artificial Intelligence

- The **theory and development of computer systems** able to perform tasks normally requiring **human intelligence**, such as visual perception, speech recognition, decision-making, and translation between languages.
- Any technique which enables computers to mimic human behavior.

■ Machine learning

A field of computer science that uses statistical techniques to give computer systems **the ability to "learn"** with data **without being explicitly programmed** and **progressively improve performance** on a specific task .

■ Pattern Classification

- Pattern classification is a sub-topic of machine learning.
- Pattern classification can be defined as a **technique to classify data** (patterns) based either on a **priori knowledge** or **statistical information** extracted from the patterns.
- Pattern recognition **automatically discover the regularities** in data through the use of learning algorithms.

Human Perception

- Humans have developed highly sophisticated skills for sensing their environment and taking actions according to what they observe, e.g.,
 - Recognizing a face,



- Understanding spoken words,
 - Reading handwriting,

504192

- Distinguishing fresh food from its smell.

- We would like to give similar capabilities to machines.

What is Pattern Classification?

Classification



Apples

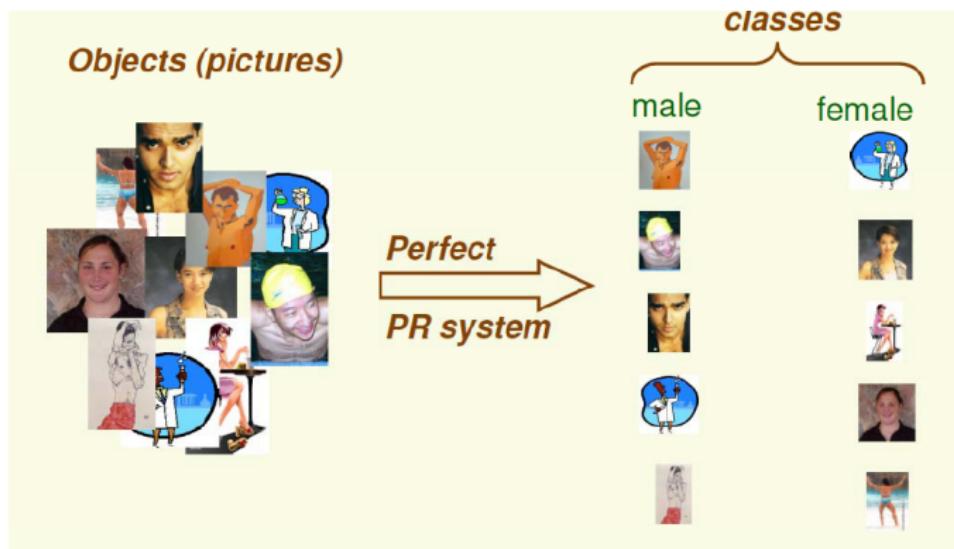


Oranges



What is Pattern Classification?

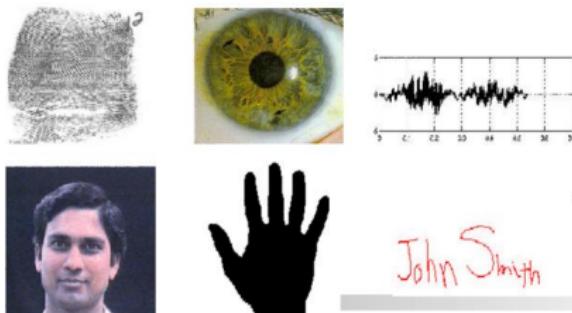
- Assign an unknown pattern to one of several known categories (or classes).



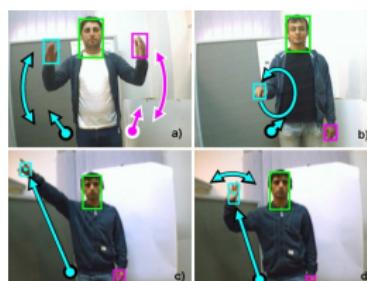
What is a Pattern?

- A pattern could be an **object** or **event**.

Biometric pattern



Hand gesture pattern



What is Pattern Classification?

- A **pattern** is an entity, vaguely defined, that could be given a name, e.g.,
 - fingerprint image,
 - handwritten word,
 - human face,
 - speech signal,
 - DNA sequence,
 - ...
- **Pattern classification** is the study of how machines can
 - observe the environment,
 - learn to distinguish patterns of interest,
 - make sound and reasonable decisions about the categories of the patterns.

Pattern Class

- A collection of “similar” objects.

Female

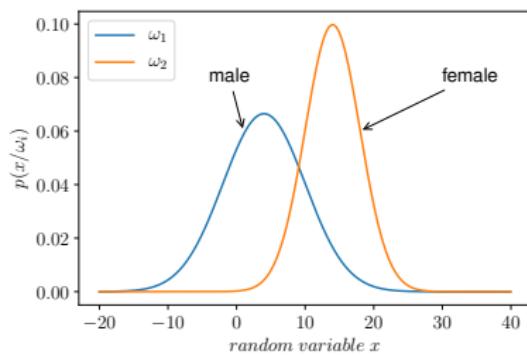
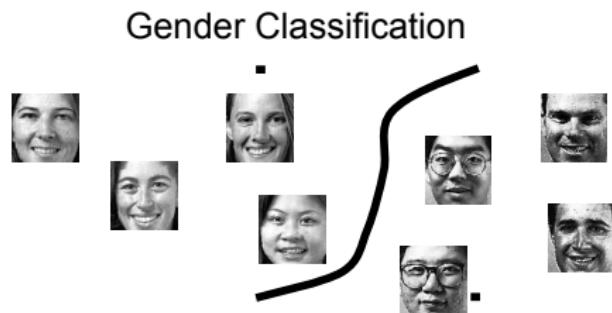


Male



How do we model a Pattern Class?

- Typically, using a statistical model.
- Probability density function (e.g., Gaussian)



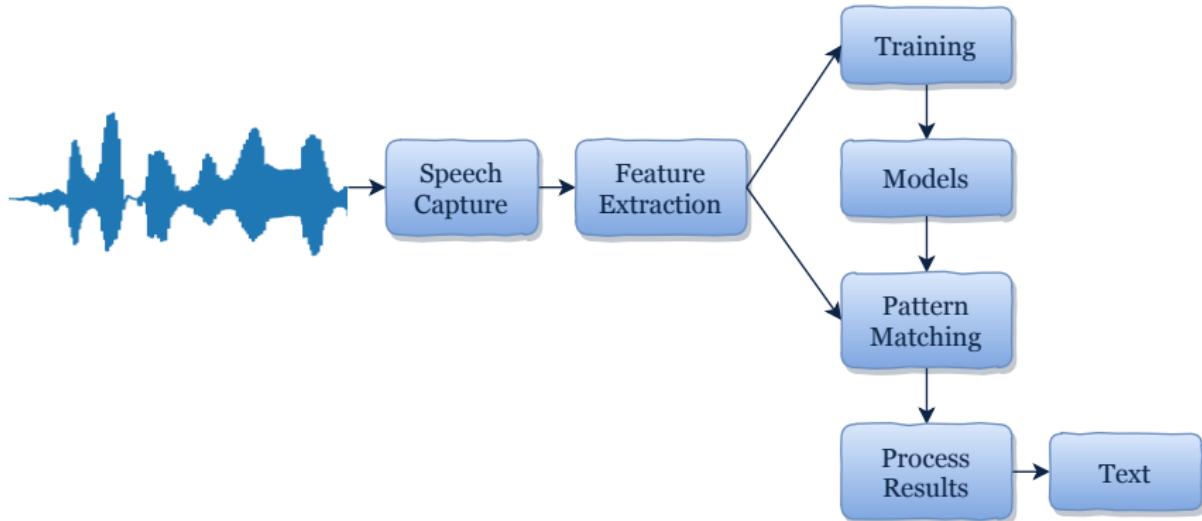
Applications

Machine Perception

Build a machine that can recognize patterns:

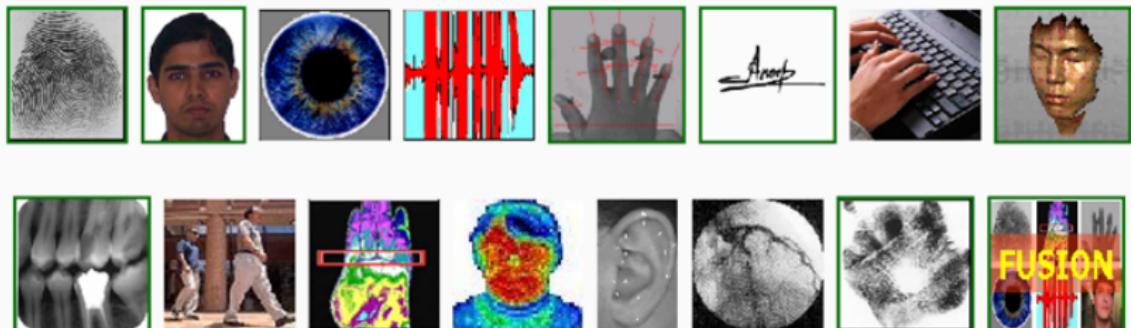
- Speech recognition
- Biometric recognition
- Fingerprint identification
- Face recognition
- OCR (Optical Character Recognition)
- DNA sequence identification
- Autonomous navigation

Applications: Speech Recognition



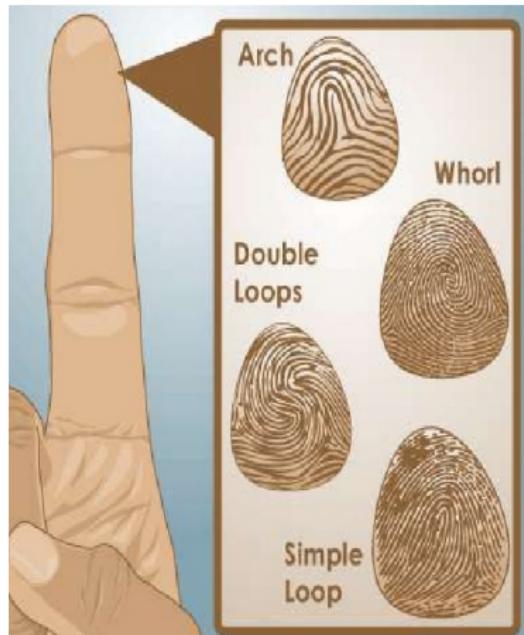
Applications: Biometric Recognition

- Fingerprint
- Eyes - Iris recognition
- Face recognition, identification/verification
- Speech recognition
- Finger geometry recognition
- Hand geometry recognition
- Signature recognition
- Eyes - Retina recognition
- Fingerprint recognition
- Typing recognition
- Gait recognition
- DNA identification



Applications: Fingerprint Identification

- No two people have the same fingerprints
- Fingerprints can solve crimes.
- Fingerprints are impressions created by ridges on the skin.
- Ridges form before a baby is born and maintain their pattern throughout life.
- As we grow, the pattern gets larger, but does not change.



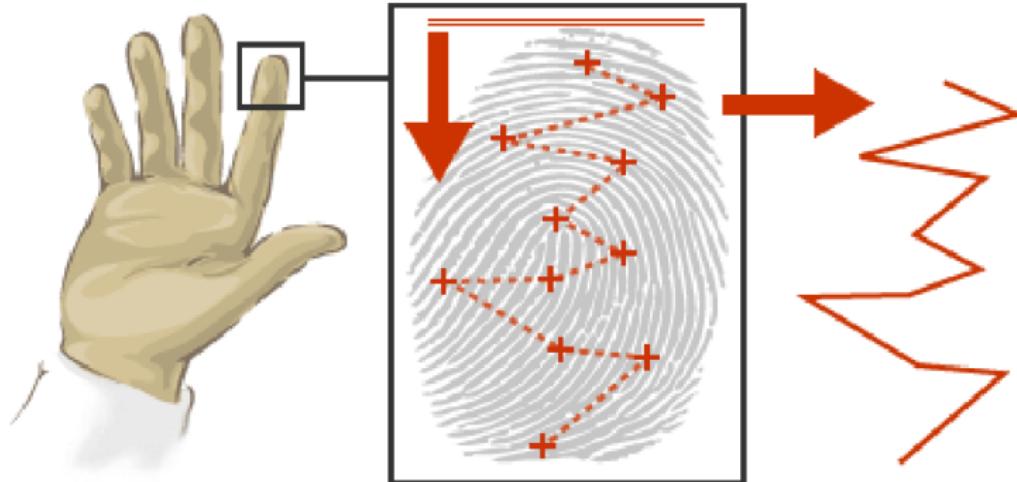
Applications: Fingerprint identification

How Fingerprint scanners record identifies:

1. Individuals index
finger is pressed
onto scanner

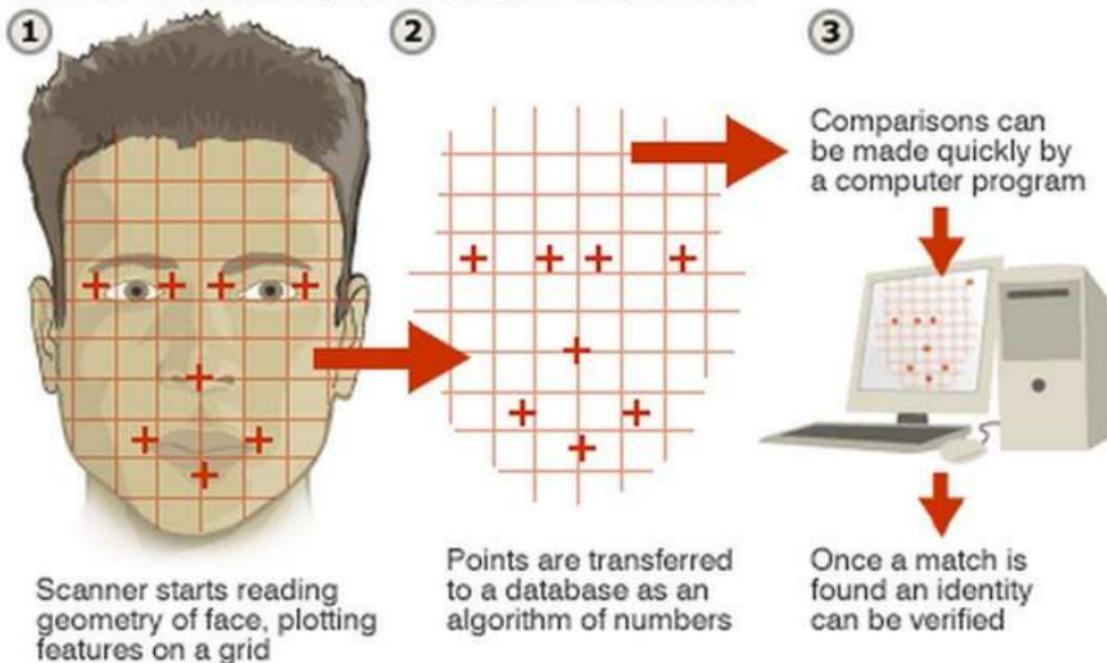
2. Scanner reads unique
pattern plotting specific
distinctions (minutiae)

3. Points are linked forming
a pattern recorded
as an algorithm used for
comparison

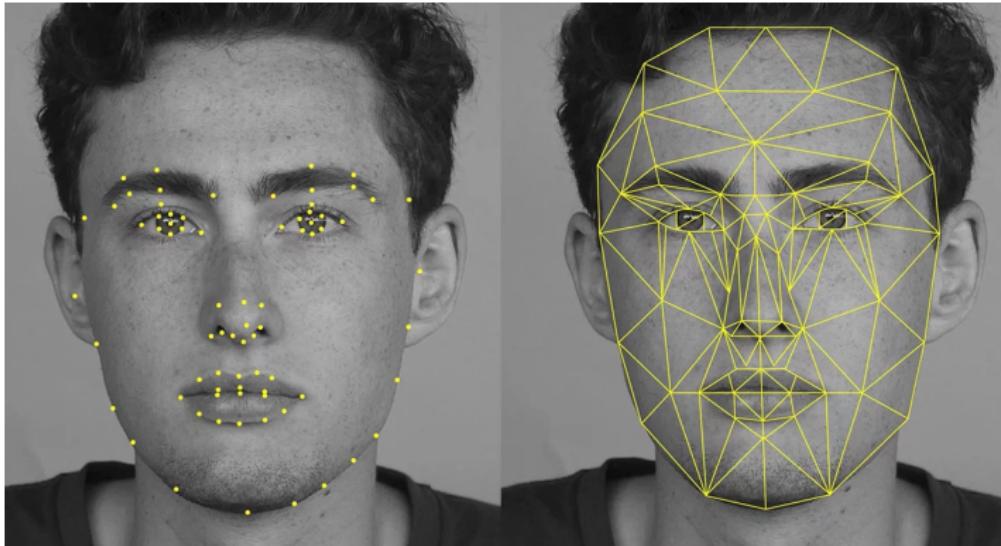


Applications: Face Recognition

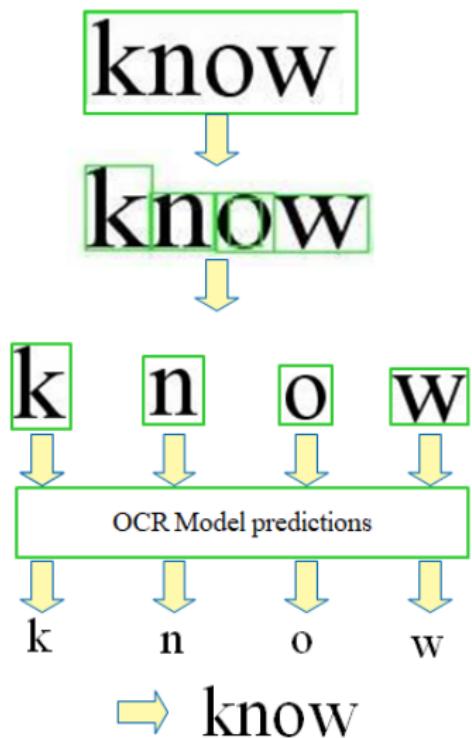
HOW 2D FACIAL SCANNERS RECORD IDENTITIES



Applications: Face Recognition



Applications: Optical Character Recognition (OCR)



1. Differentiate word contours associated with Image.
2. Differentiate letter contours associated with word contours associated with word contour image.
3. Preprocess letter images according to trained OCR input
4. Consolidate predictions associated OCR model to text.

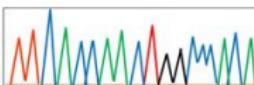
Applications: DNA Identification

***Obtain DNA-bearing
tissue samples:
blood, saliva, hair
follicles, feathers,
scales***



***Extract DNA
from Cells***

Sequence DNA



...TTCAACAAACAGGCCAC...

***Identify
matching DNA
sequence(s) in
databases***

TTCAACAAACAGGCCAC
TTCACCAACAGGCCAC
TTCATCTACAGCCCCAC

Applications: Autonomous Navigation



Applications: Cancer Detection

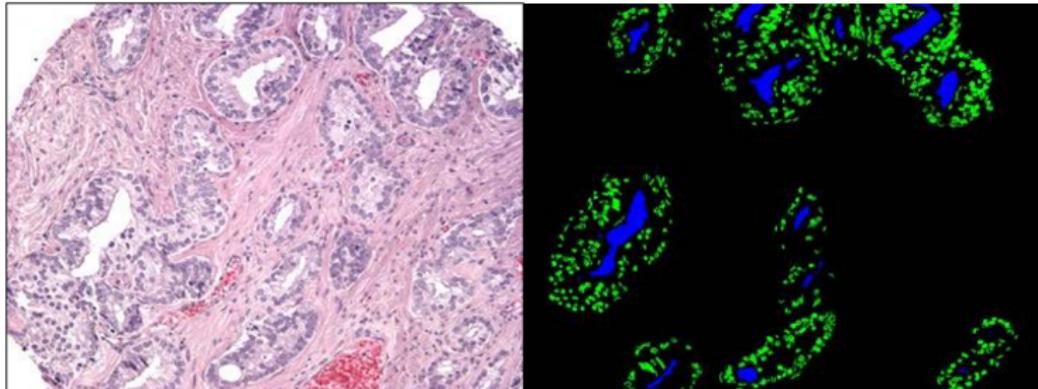


Figure: Cancer detection and grading using microscopic tissue data

Applications: Land Cover Detection



Figure: Land cover classification using satellite image

Applications: License Plate Recognition

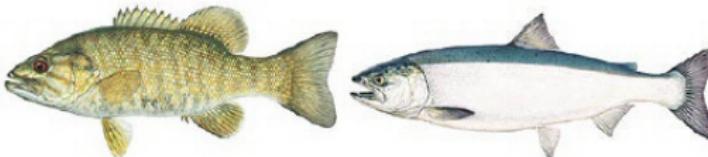


Figure: License plate recognition: US license plates

A classic example to understand Pattern Classification

An Example

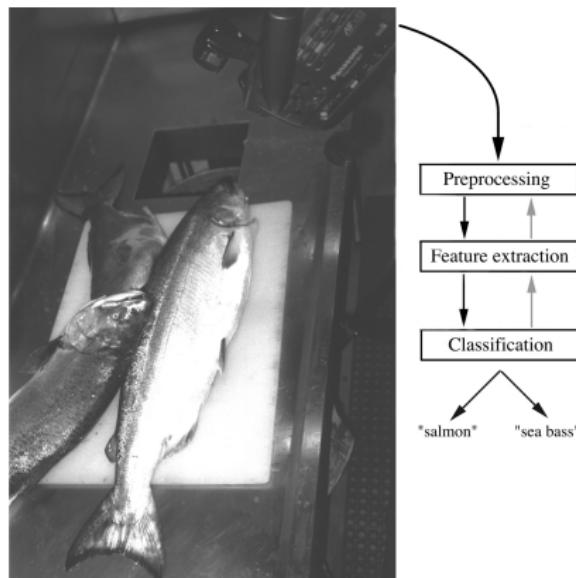
- “Sorting incoming Fish on a conveyor according to species using optical sensing”
- Species
 - Sea bass
 - Salmon



- Set up a camera and take some sample image to extract features
 - Length
 - Lightness
 - Width
 - Number and shape of fins
 - Position of the mouth, etc.
- This is the set of all suggested features to explore for use in our classifier.

An Example

- What can cause problems during sensing?
 - **lighting conditions**,
 - position of fish on the conveyor belt,
 - camera noise, etc.
- Use a segmentation operation to **isolate fishes** from one another and from the **background**.
- Information from a single fish is sent to a **feature extractor** whose purpose is to reduce the data by measuring certain features.
- The features are passed to a classifier.



An Example: Classification

- Select the length of the fish as a possible feature for discrimination

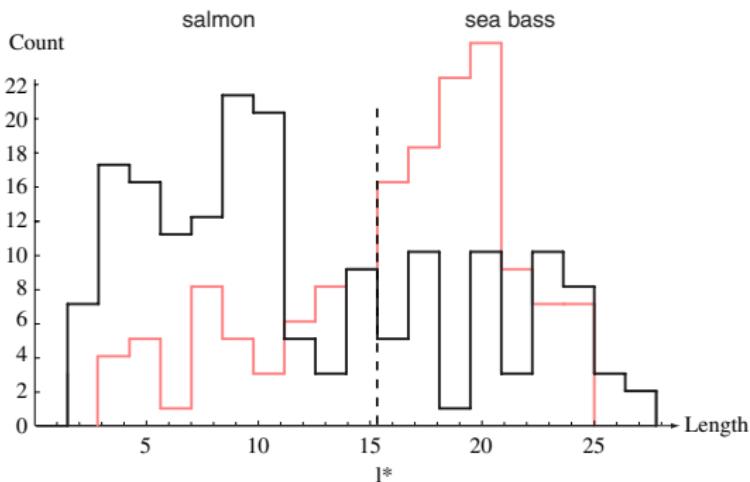


Figure: Histograms for the length feature for the two categories. No single threshold value l^* (decision boundary) will serve to unambiguously discriminate between the two categories; using length alone, we will have some errors. The value l^* marked will lead to the smallest number of errors, on average.

An Example: Classification

- The length is a poor feature alone
- Select the lightness as a possible feature

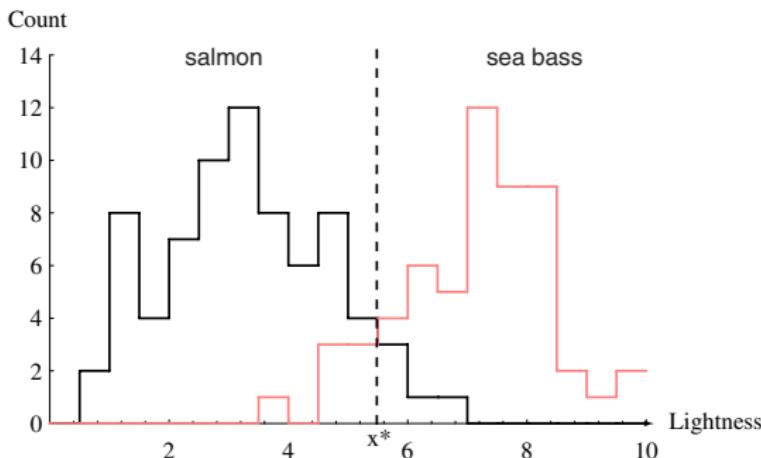


Figure: Histograms for the lightness feature for the two categories. No single threshold value x^* (decision boundary) will serve to unambiguously discriminate between the two categories; using lightness alone, we will have some errors. The value x^* marked will lead to the smallest number of errors, on average.

Threshold decision boundary and cost relationship

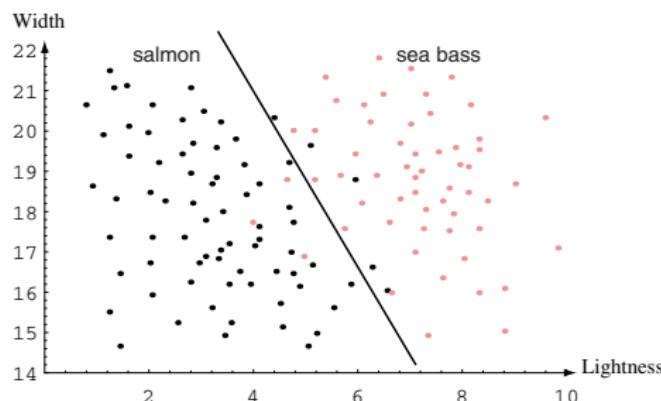
- Move our decision boundary toward smaller values of lightness in order to *minimize the cost* (reduce the number of sea bass that are classified salmon)



Task of decision theory

An Example: Feature vector

- Adopt the lightness and add the width of the fish
- We can use two features in our decision:
 - lightness: x_1
 - width: x_2
- Each fish image is now represented as a point (**feature vector**) x in two-dimensional **feature space**.



$$x = [x_1 \quad x_2]^T$$
$$\downarrow \quad \downarrow$$

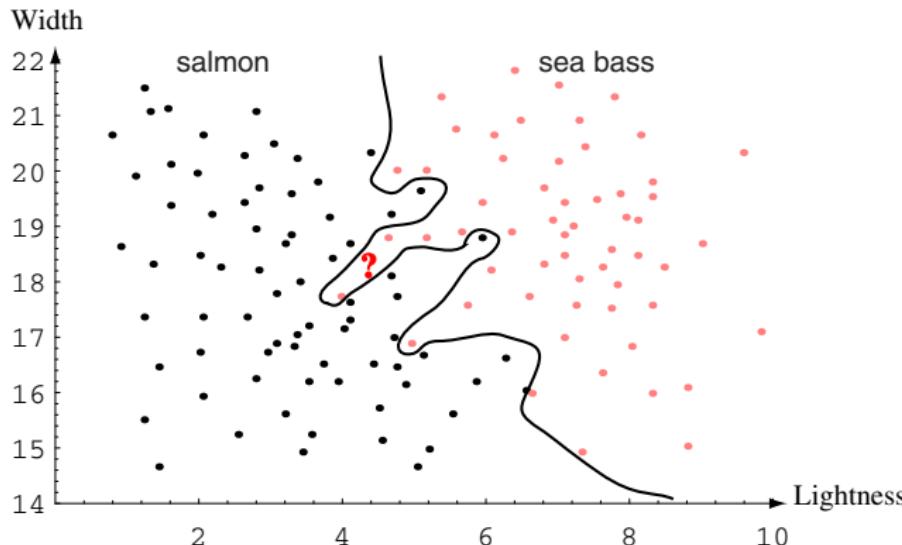
Lightness Width

An Example: Feature vector

- We might add other features that are not correlated with the ones we already have. A precaution should be taken not to reduce the performance by adding such “noisy features”.
- Does adding more features always improve the results?
 - unreliable features.
 - Be careful about correlations with existing features.
 - Be careful about measurement costs.
 - Be careful about noise in the measurements.
- Is there some *curse* for working in very high dimensions?

An Example: Feature vector

- Ideally, the best decision boundary should be the one which provides an optimal performance.



An Example: Issue of generalization

- How can we manage the *tradeoff* between complexity of decision rules and their performance to unknown samples?
- Our satisfaction is premature because the central aim of designing a classifier is to correctly classify novel input



Issue of generalization

An Example: Decision Boundary

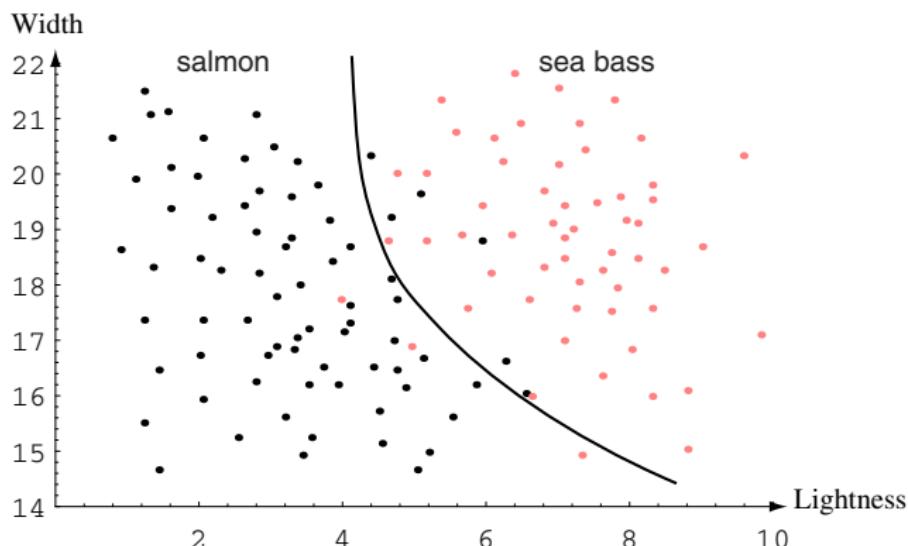


Figure: The decision boundary shown might represent the optimal trade off between performance on the training set and simplicity of classifier.

More on Complexity

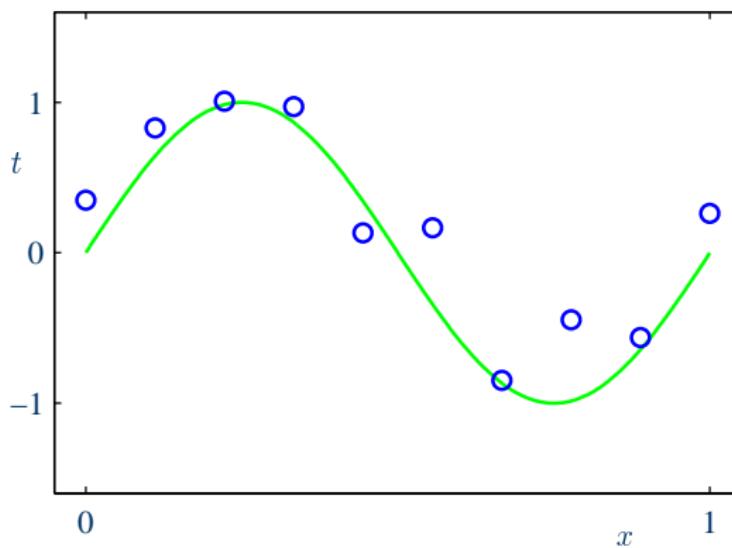
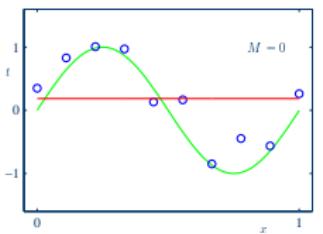
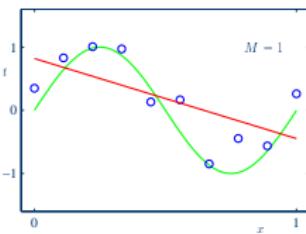


Figure: Regression example: plot of 10 sample points for the input variable x along with the corresponding target variable t . Green curve is the true function that generated the data.

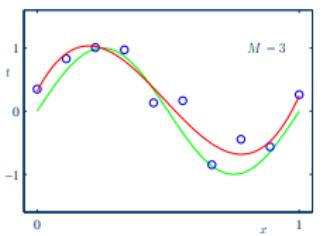
More on Complexity



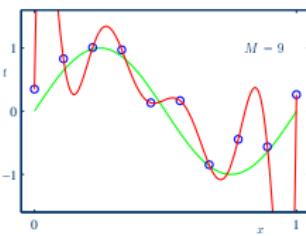
(a) 0'th order polynomial



(b) 1'st order polynomial



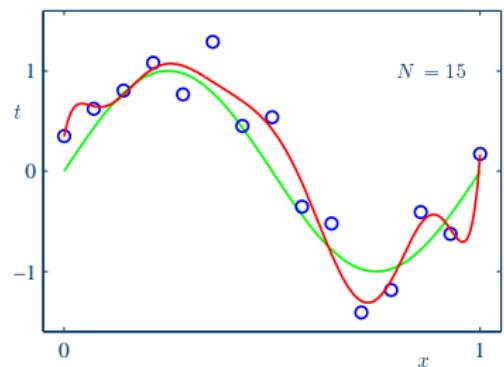
(c) 3'rd order polynomial



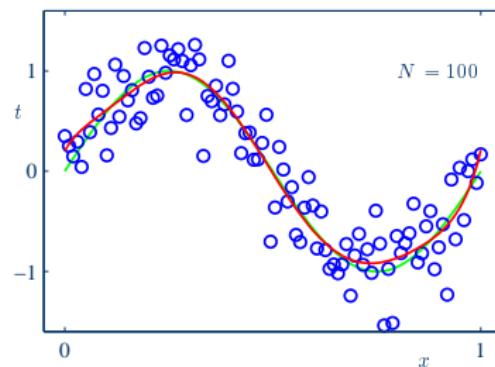
(d) 9'th order polynomial

Figure: Polynomial curve fitting: plots of polynomials having various orders, shown as red curves, fitted to the set of 10 sample points

More on Complexity



(a) 15 sample points



(b) 100 sample points

Figure: Polynomial curve fitting: plots of 9'th order polynomials fitted to 15 and 100 sample points.

Pattern recognition system and design cycle

Pattern Recognition Models

There are three main models of pattern recognition:

- **Statistical:**

- to identify where specific piece belongs (for example, whether it is a cake or not).
 - use of statistics to learn from examples.

- **Syntactic/Structural:**

- to define a more complex relationship between elements taking into account more complex interrelationships between attributes.
 - Looks at clear structure in the patterns.
 - An example of this would be diagnosis of the heart with ECG measurements.

- **Template Matching:**

- To match the object's features with the predefined template and identify the object by proxy.
 - One of the uses of such model is plagiarism checking.

Pattern Recognition Systems

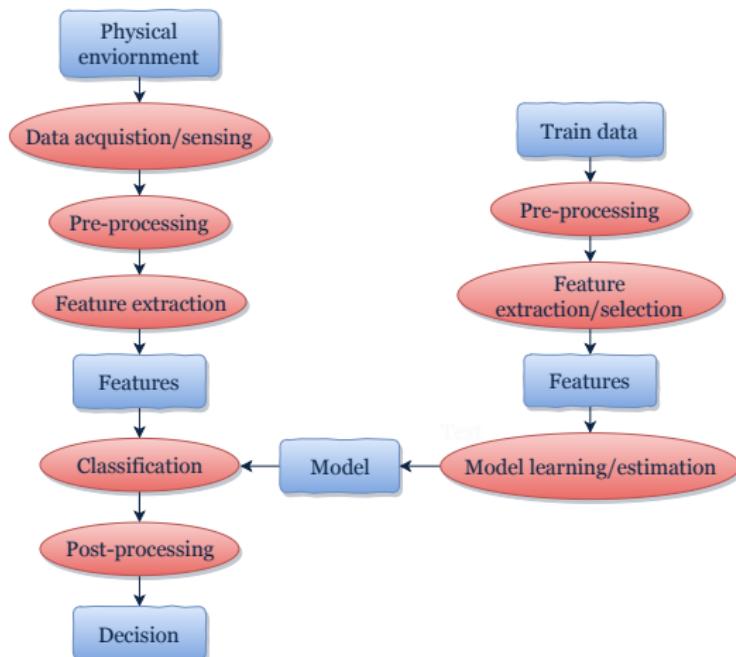


Figure: Object/Process diagram of a pattern recognition system

Pattern Recognition Systems

■ Data acquisition and sensing:

- Use of transducer (camera or microphone)
- Measurements of physical variables.
- Important issues: bandwidth, resolution, sensitivity, distortion, SNR, latency, etc.

■ Pre-processing:

- Removal of noise in data.
- Isolation of patterns of interest from the background.
- Segmentation and grouping

■ Feature extraction:

- Finding a new representation in terms of features.
- Features should be well separated and should not overlap (Discriminative features)
- Invariant features with respect to translation, rotation and scale.
- Depends on the characteristics of the problem domain. Simple to extract, invariant to irrelevant transformation insensitive to noise.

Pattern Recognition Systems

■ Model learning and estimation:

- Learning a mapping between features and pattern groups and categories.
- Unsatisfied with the performance of classifier and want to jump to another class of model

■ Classification:

- Using features and learned models to assign a pattern to a category.

■ Post-processing:

- Evaluation of confidence in decisions.
- Exploitation of context to improve performance.
- Combination of experts.

The Design Cycle

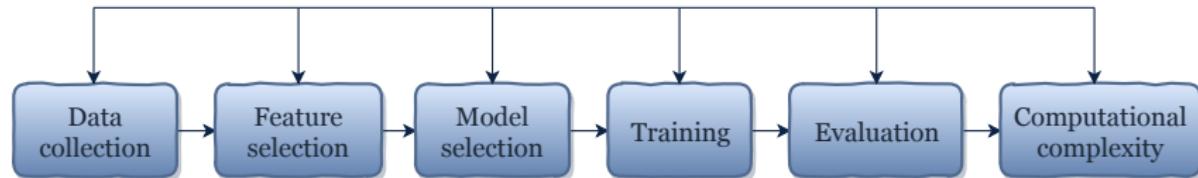


Figure: Design cycle

■ Data Collection

- Collecting training and testing data.
- How do we know when we have collected an adequately large and representative set of examples for training and testing the system?

The Design Cycle

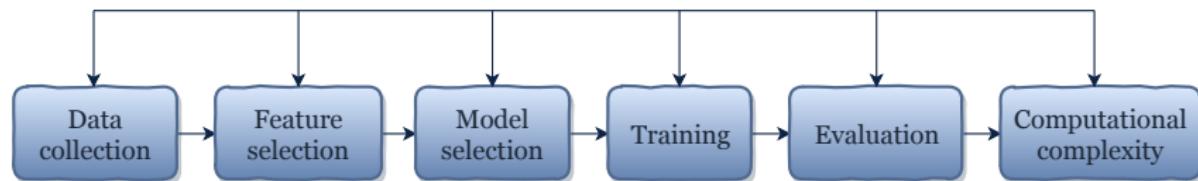


Figure: Design cycle

■ Feature Selection

- Domain dependence and prior information.
- Computational cost and feasibility.
- Discriminative features.
 - Similar values for similar patterns.
 - Different values for different patterns.
- Invariant features with respect to translation, rotation and scale.
- Robust features with respect to occlusion, distortion, deformation, and variations in environment.

The Design Cycle

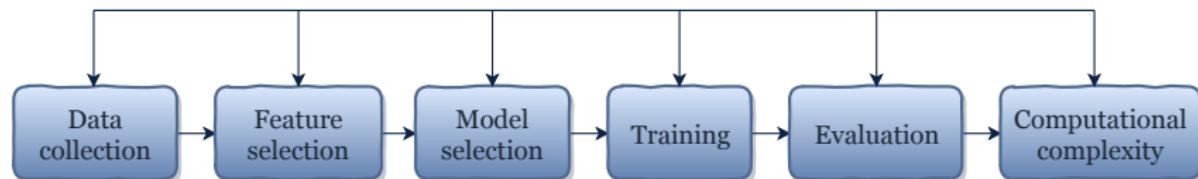


Figure: Design cycle

■ Model selection

- Domain dependence and prior information.
- Definition of design criteria.
- Parametric vs. non-parametric models.
- Handling of missing features.
- Computational complexity.
- Types of models: templates, decision-theoretic or statistical, syntactic or structural, neural, and hybrid.
- How can we know how close we are to the true model underlying the patterns?

The Design Cycle

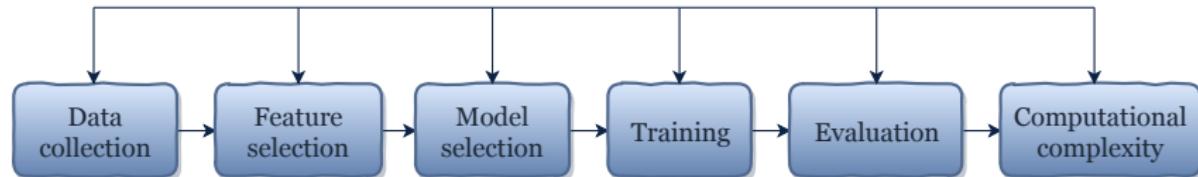


Figure: Design cycle

■ Training

- How can we learn the rule from data?
- Supervised learning: a teacher provides a category label or cost for each pattern in the training set.
- Unsupervised learning: the system forms clusters or natural groupings of the input patterns.
- Reinforcement learning: no desired category is given but the teacher provides feedback to the system such as the decision is right or wrong.

The Design Cycle

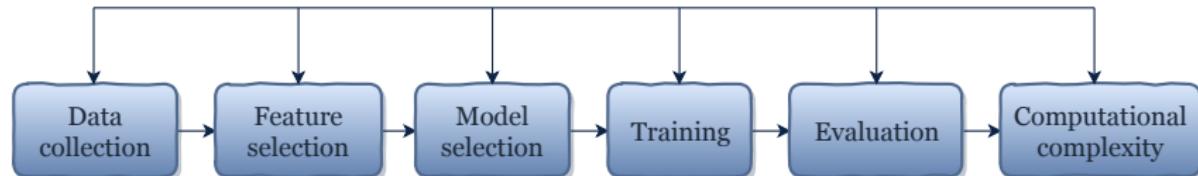


Figure: Design cycle

■ Evaluation

- How can we estimate the performance with training samples?
- How can we predict the performance with future data?
- Problems of overfitting and generalization.

■ Computational Complexity

- What is the trade-off between computational ease and performance?
- How an algorithm scales as a function of the number of features, patterns or categories?

Different Learning Approaches

■ Supervised learning/Classification

- A teacher provides a category label or cost for each pattern in the training set

■ Unsupervised learning/Clustering

- The system forms clusters or "natural groupings" of the input pattern (no explicit teacher)

■ Semi-supervised learning

- Semi-supervised learning is the problem of learning from examples for which you have labels for only a (small) subset.

■ Reinforcement learning

- Learning with critic, no desired category is known; instead, the only teaching feedback is that the tentative category is right or wrong. It utilizes reward function to learn. Ex: Autonomous driving

References

- [1] Omer Cengiz CELEBI. *Neural Networks and Pattern Recognition* [Online, Last accessed on 26 Dec 2019].
https://www.byclb.com/TR/Tutorials/neural_networks/Index.aspx.
- [2] Richard O Duda, Peter E Hart, and David G Stork. *Pattern classification*. John Wiley & Sons, 2012.



Thank you!